

Automatic Parallelisation of Rust Programs at Compile Time

Project Proposal

TODO: Write in proper sentences more objectively

1 Problem

Kish (2002) claims Moore's Law is dying/dead and processors have been gaining more and more cores (Geer 2005). Parallelising programs to take advantage of the additional cores has some difficulty and often requires the programmer to make significant changes to the source code. It would be much nicer if the programmer could write normal sequential code, and the parts that can be run in parallel are automatically converted into parallel code. Also existing sequential code could take advantage of multithreading without rewriting.

2 Approach

People have automated parallelisation of sequential FORTRAN and C code in the past (D'Hollander, Zhang and Wang 1998; Baskaran, Ramanujam and Sadayappan 2010). Instead of trying to parallelise one of these languages, I will try to automatically parallelise sequential rust programs, as the rust language has guaranteed memory safety and threads without data races (*The Rust Programming Language* 2017).

I am going to write a compiler plugin which will analyse the abstract syntax tree of another rust program. Any parts of the abstract syntax tree that could be run in parallel would be run in it's own thread instead.

2.1 Requirements

- A computer to program on
- Source code of some rust programs to parallelise
- Rust compiler
- Atom

2.2 Timeline

TODO: Add estimated dates

Below is an estimated timeline of the project.

- Write this proposal
- Submit the ethical review

- Create two plugins, a linter and a syntax extension
 - Get the linter plugin to:
 - * Analyse the programs statements to see what variables each statement depends on and what variables the statement modifies (if any)
 - * Produce a dependency tree for the entire program based on this analysis
 - * Look for areas in the tree which do not depend on one another, these areas could be run in parallel
 - * Estimate the speed of each statement for the dependency tree
 - * Record the areas that could be changed into a file
 - Get the syntax extension plugin to:
 - * Read the file the Linter plugin creates
 - * First try one statement in parallel to test it works
 - * Then try to run everything in parallel to test it works
 - * Then only run parts in parallel if it would be faster to run in parallel (maybe compile to if $n < 1000$ then serial else parallel)
- Do some testing
- Write a report

2.3 Possible Extensions

If I'm ahead of schedule with the coding section of this project, I could look into parallelising 'if' statements which have a slow condition. Each branch of the 'if' would be run in a separate thread using cloned data. When the condition is finally calculated the incorrect branches would need to be deleted. This kind of parallelisation is different from the project plan as some threads are "thrown away".

Another possible extension upon the previous extension is to utilise the GPU using CUDA in cases where a large number of threads are doing the exact same task on different data. I feel it is very unlikely that I would have time for this extension, and I'm unsure of how much real world code would be written in such a way that it could be automatically run on the GPU efficiently.

3 Evaluation

3.1 Disadvantages to the Chosen Approach

While I think the approach I have chosen is the best option to solve the problem, there are a few downsides to this proposed approach.

To take an existing sequential rust program and compile it requires a few steps. The compiler plugin crate would need to be imported, and every function would need to be annotated so that the plugin has edit access to the abstract syntax tree for that function.

Due to the limitation of rust compiler plugins, the program would need to be compiled twice. The first compile would allow the entire syntax tree to be analysed for parallelisable sections. The second compile would edit the syntax tree to move the parallelisable sections into separate threads.

Project Risks:

- Compiler code too complicated to use
- Size of task is too large
- No idea how to analyse the speed to statements at compile time yet
- “Independent tasks can be run in parallel” is true in my head, but maybe not in practice
- Rust may not give me all the guarantees that I think it would give me

3.2 Other Approaches

There are other approaches to the problem that I am trying to solve (some of which have already been implemented). I will explain these other approaches and why I didn't choose them.

3.2.1 Make a custom rust compiler

One of the downsides of the chosen approach is that the program has to be compiled twice. I may be able to bypass this limitation by making a fork of the rust compiler instead of using compiler plugins. It would then be possible to examine the entire abstract syntax tree and then edit it afterwards. However, this approach would add significant complexity to the project as the rust compiler code is very complex.

3.2.2 Use a different language from rust

There are many programs out there, written in many different languages. So why rust? Rust is not even in the top 25 of most popular programming languages according to StackOverflow (2017). Using a more popular language would make this tool more useful for more programs. Rust isn't unpopular as it is the most loved language and the 10th language on the most want to learn list (StackOverflow 2017). The main reason for using rust is the guaranteed memory safety and threads without race conditions which would help significantly with the task.

3.2.3 Manually annotate the parts of the program that are parallelisable

(Dagum and Menon 1998)

3.2.4 Just run the sequential code

Running code in multiple threads has some overhead.

3.3 Measuring Project Success

I can evaluate the successfulness of my solution by measuring the speedup of parallelising the program vs the original serial code. I'll look at existing programs written in rust to see any real world impact.

4 References

- Baskaran, Muthu, Jj Ramanujam and P Sadayappan (2010). "Automatic C-to-CUDA code generation for affine programs". In: *Compiler Construction*. Springer, pp. 244–263.
- Dagum, Leonardo and Ramesh Menon (1998). "OpenMP: an industry standard API for shared-memory programming". In: *IEEE computational science and engineering* 5.1, pp. 46–55.
- D'Hollander, Erik H, Fubo Zhang and Qi Wang (1998). "The fortran parallel transformer and its programming environment". In: *Information sciences* 106.3-4, pp. 293–317.
- Geer, David (2005). "Chip makers turn to multicore processors". In: *Computer* 38.5, pp. 11–13.
- Kish, Laszlo B (2002). "End of Moore's law: thermal (noise) death of integration in micro and nano electronics". In: *Physics Letters A* 305.3, pp. 144–149.
- StackOverflow (2017). *Stack Overflow Developer Survey 2017*. URL: <https://insights.stackoverflow.com/survey/2017> (visited on 09/10/2017).
- The Rust Programming Language* (2017). URL: <https://www.rust-lang.org/en-US/> (visited on 09/10/2017).